

Domain

Conspiracy theories.

I'm somewhat familiar with the topic. (Beyond reading the news), I've read a bit of the academic literature on the subject of American conspiracy theories, so I'm familiar with the basics of the major conspiratorial discourses, their contours, logics, subject-matter, etc. My background in epistemology could help, and I did a semester-long fellowship on the epistemology of bullshitting, which is not the same thing, but is one other kind of logically deviant discourse.

Data

I've already collected a little over 200 primary source texts that represent a decent cross-section of what you might call "the canon" or most significant/influential conspiratorial works. They're mostly American, mostly 20th Century. My selection process was to go line-by-line through the bibliography of academic book(s) on conspiracy theories and wikipedia, and download as many of the primary source works as I could. Most of the 200 sources are book-length, some article length. Almost all the ones that aren't in txt format I can convert to txt. I've got the bulk of the most important works, but if I can get more texts and you think that I should, then I will try to find more.

For my MVP, I'd like to perform topic modeling on the texts and cluster them into subject-matters (i.e. aliens, new world order, reptiles, etc.) My hope is that some unexpected topics or clusters emerge, and/or that I'll discover something interesting about topic overlap across clusters.

If there is time, I'd also like to do the same thing, but for style (probably by knocking out topic words) and for purported epistemic structure (this will be tricky, but maybe a combination of sentence structure and keywords like {'rationalism':('therefore','because'),..., 'empirical': 'saw','witnessed','heard')... 'authority':('according to','wrote','discovered'...}) could yield some interesting results. The subject, style, and epistemic clusterings will be distinct, but where they overlap could be interesting.

Ideally, I'd like to create some kind of model that would allow me input an unseen text and output how close a fit it is with the different clusters. (If necessary for supervised training, I could use texts that I label as non-conspiratorial which use similar vocabularies to the clusters.) A flask app could be really fun.

Another stretch goal would be to construct a graph of Twitter networks, and create a snapshot of how conspiratorial messages proliferate (probably with users as vertices, and retweets representing paths), and see how much clustered structure there is, whether clusters interact, etc.

Known unknowns

How long any of these steps will take, whether stylistic and epistemic detection is even possible, and whether I need more data. I'm familiar with various graph search algorithms, but haven't done graph discovery or used the Twitter API.

